

NLP - Fake news detection

Simbiat Musa Georg F.K. Höhn

Ironhack

19 June 2025

Project overview

Given

- ▶ dataset of headlines annotated as fake (0) or real (1)

Aim

- ▶ classify news headlines in unseen data (using machine learning)

Deliverables

- ▶ Python pipeline (+ csv with model training results)
- ▶ csv with predictions for test set

Data overview

Methodology

Training results

Conclusion

Data overview

- ▶ 34152 rows of annotated data
 - ▶ fake news: 17572
 - ▶ real news: 16580
- ▶ relatively balanced

Word cloud for data annotated as fake

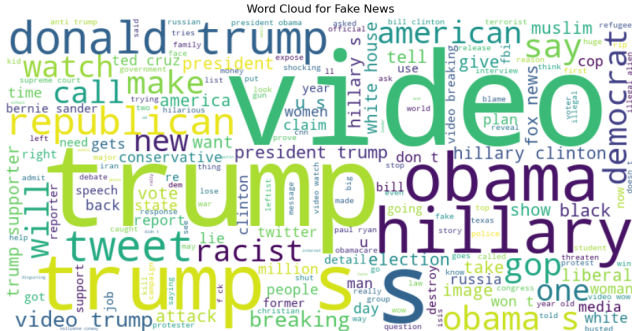


Figure 2: Word cloud for fake headlines

Methodology

Structure

- ▶ one main .py file for data exploration, model experimentation and output generation
- ▶ helper.py with functions for
 - ▶ cleaning strings
 - ▶ generate, print and save model evaluations (to csv)
 - ▶ removing stop words (not needed with TF-IDF vectoriser)
 - ▶ lemmatizer (currently not used)
 - ▶ huggingface pipeline for transformer models

General pipeline

- ▶ EDA
- ▶ clean data (generate new columns)
- ▶ train-test split (20% test size)
- ▶ vectorise train and test sets
- ▶ train (and tune?) models
- ▶ compare based on test accuracy
- ▶ run best model(s) on target data and save
- ▶ (run best models of different types on target data and calculate inter-annotator agreement)

Vectorisation

- ▶ TF-IDF
- ▶ GloVE (glove-wiki-gigaword-100)

ML-algorithms

- ▶ Logistic Regression
(`sklearn.linear_model.LinearRegression`)
- ▶ Random Forest
(`sklearn.ensemble.RandomForestClassifier`)
- ▶ KNN (`sklearn.neighbors.KNeighborsClassifier`)
- ▶ XGBoost (`xgboost.XGBClassifier`)

Regarding transformer models

- ▶ tried pipelines with
 - 1) jy46604790/Fake-News-Bert-Detect
 - 2) omykhailiv/bert-fake-news-recognition
- ▶ both performed abysmally: everything is fake, see result for 1
- ▶ over-sensitive? issue with our data pre-processing?




Figure 3: jy46604790/Fake-News-Bert-Detect

Training results

model-id	params	acc-train	accuracy
logreg-final	miter=500	0.917	0.908
logreg-1000	miter=1000	0.917	0.908
xgb-final	est=500,mdepth=100,lr=0.3, α =0.1	0.990	0.906
xgb	est=500,mdepth=100,lr=0.3	0.990	0.905
xgb	est=400,mdepth=100,lr=0.5	0.990	0.905
rndforest-final	est=300,min_samp_leaf=2	0.938	0.895
xgb	est=100,mdepth=50,lr=0.04	0.915	0.882
rndforest	est=100	1.000	0.878
logreg-glove	miter=1000	0.869	0.870
xgb	est=200,mdepth=50,lr=0.005	0.877	0.850
rndforest	est=300,mdepth=30	0.868	0.842
xgb	est=10,mdepth=50,lr=0.04	0.861	0.837
knn	k=3	0.915	0.817
knn	k=5	0.883	0.804
knn	k=10	0.816	0.780
omykhailiv	defaults	0.514	0.517
jy46604790	defaults	0.514	0.517

- ▶ best performing: logistic regression model
 - ▶ max-iterations did not seem to make a difference
- ▶ xgb and RandomForest very close by
 - ▶ but: high risk of overfitting
 - ▶ longer training times/higher complexity

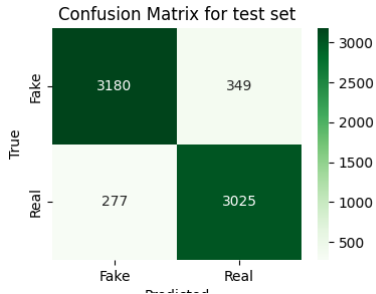


Figure 4: Confusion matrix for logistic regression model

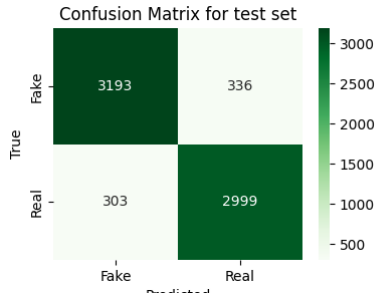


Figure 5: Confusion matrix for CGB model

Conclusion

- ▶ logistic regression model offers best performance

Collaboration

- ▶ using py files to avoid notebook consistency issues with git
- ▶ # %%: useful VS Code option for generating jupyter-like cells

Outlook/reflection

- ▶ further experimentation with vectoriser settings
- ▶ actually implement lemmatisation (oops)
- ▶ some confusion in our raw performance results.csv
 - ▶ some RandomForest models seemed to perform better than final model, but probably due to earlier mistakes in preprocessing?
 - ▶ lesson: also note changes to preprocessing or cleanly reset logging files

Thanks for your attention!